

Backend engineer and AI evaluation specialist. Most recently built backend safety-scoring systems at Oasis Health; previously shipped adversarial testing, failure-mode taxonomies, and structured-evaluation pipelines at Scale AI and Persona Identities. I bring evaluation discipline to production code. Open to backend and AI-evaluation roles.

EXPERIENCE

- Oasis Health · Backend Engineer

Jan 2026 to May 2026

 - Built the backend safety-scoring service that ranks consumer products against verified contaminant thresholds and surfaces risk indicators to the mobile clients. **500k+** requests per day across **20+** managed endpoints.
 - Designed the data layer behind the scoring engine and rolled out the new scoring pipeline to **100%** of production traffic.
 - Hardened the production stack (auth, rate limiting, observability) so mobile clients can be trusted in front of real consumer purchase decisions.
- Persona Identities · Senior Product Development Specialist

Oct 2025 to Nov 2025

 - Delivered two end-to-end LLM evaluation projects for enterprise clients, including pipeline review, reporting, and post-mortem documentation.
 - Implemented security fixes that mitigated incoming risk across **5+** high-priority enterprise clients, closing identity and reasoning vulnerabilities before deployment.
 - Authored evaluation reports and failure-mode taxonomies that fed into the team's debugging and refinement work.
- Scale AI · Operations Associate, Trust and Safety (Eval Engineer)

Apr 2024 to Jul 2025

 - Built adversarial prompts and structured evaluations that surfaced harmful, biased, and unreliable outputs for **5+** enterprise clients across **20+** shipped projects.
 - Authored weekly reports analyzing drift, misalignment, hallucinations, and behavioral risk using SQL and Python.
 - Designed evaluation criteria, rubrics, and taxonomies that improved reviewer precision by **30%+** and lifted reproducibility across high-volume projects.
- Scale AI (via Hire Art) · Generative AI Specialist → Queue Manager

Jan 2024 to Apr 2024

 - Evaluated LLM outputs for safety, bias, hallucinations, and factual accuracy; ranked completions against gold standards.
 - Diagnosed pipeline breakdowns and quantified reviewer throughput and variance with Python and SQL.
- Meta · Data Labeling Analyst, Redteaming

Nov 2023 to Dec 2023

 - Performed adversarial testing on production LLMs and delivered detailed failure-mode reports that fed into model robustness work.
- LinkedIn · Engineering Linguist, Seeker Relevance AI

Jun 2022 to Jul 2023

 - Boosted ML model accuracy by **30%** on Seeker Relevance through targeted linguistic annotation and preprocessing improvements.
 - Authored annotation guidelines for the Seeker Relevance training corpus and partnered with ML engineering to convert linguistic findings into ranking and recall lifts.

EDUCATION

UC Davis	2022
Data Analytics / Data Science Bootcamp	
UC Santa Cruz	2021
B.A. Linguistics	

PROJECTS

- Mental Health Risk Predictor**, ML model identifying income inequality as the strongest predictor of community mental-health risk. [Case study](#) · [Code](#)
- Citi Bike Tableau**, live NYC trip dashboard on Tableau Public. [Case study](#) · [Live](#)
- Full portfolio at jonahrahn.github.io.

SKILLS

BACKEND	API design · REST · database design · data pipelines · observability · auth
LLM EVAL	Adversarial prompting · failure-mode taxonomy · red-teaming · rubric design
LANGUAGES	Python (pandas, NumPy, Jupyter) · SQL · JavaScript
DATA	BigQuery · Looker · Tableau · ETL · A/B testing · data warehousing
ML / AI	PyTorch · TensorFlow · LangChain · prompt engineering · NLP
LINGUISTICS	Syntax · computational linguistics · annotation taxonomy · multilingual evaluation
INFRA	GCP · cloud infrastructure · MongoDB · PostgreSQL · Flask · React