

Backend engineer with hybrid expertise in LLM evaluation and computational linguistics.
Currently building backend systems at Oasis Health; previously shipped adversarial testing,
failure-mode taxonomies, and structured-evaluation pipelines at Persona Identities and Scale AI.
I bring evaluation discipline to production code.

EXPERIENCE

- Oasis Health · Backend Engineer

Jan 2026 – Present

 - Shipped backend APIs for Oasis's in-store product-scan flow, translating raw lab data into consumer-facing contaminant safety scores. [Add: req/day or P95 latency.]
 - Designed and owned the data-ingestion pipeline backing Oasis's contaminant database, supporting [N] product categories and [N] lab sources.
 - Hardened the production stack — auth, rate limiting, observability — so mobile clients can be trusted in front of real consumer purchase decisions.
- Persona Identities · Senior Product Development Specialist

Oct 2025 – Nov 2025

 - Delivered two end-to-end LLM evaluation projects for enterprise clients, including pipeline review, reporting, and post-mortem documentation.
 - Implemented security fixes that mitigated incoming risk across 5+ high-priority enterprise clients, closing identity and reasoning vulnerabilities before deployment.
 - Authored evaluation reports and failure-mode taxonomies that fed into the team's debugging and refinement work.
- Scale AI · Operations Associate, Trust and Safety (Eval Engineer)

Apr 2024 – Jul 2025

 - Built adversarial prompts and structured evaluations that surfaced harmful, biased, and unreliable outputs for enterprise clients.
 - Authored weekly reports analyzing drift, misalignment, hallucinations, and behavioral risk using SQL and Python.
 - Designed evaluation criteria, rubrics, and taxonomies that improved reviewer precision and reproducibility across high-volume projects.
- Scale AI (via Hire Art) · Generative AI Specialist → Queue Manager

Jan 2024 – Apr 2024

 - Evaluated LLM outputs for safety, bias, hallucinations, and factual accuracy; ranked completions against gold standards.
 - Diagnosed pipeline breakdowns and quantified reviewer throughput and variance with Python and SQL.
- Meta (via Magnit) · Data Labeling Analyst, Redteaming

Nov 2023 – Dec 2023

 - Performed adversarial testing on production LLMs and delivered detailed failure-mode reports that fed into model robustness work.
- LinkedIn · Engineering Linguist, Seeker Relevance AI

Jun 2022 – Jul 2023

 - Boosted ML model accuracy by 30% on Seeker Relevance through targeted linguistic annotation and preprocessing improvements.
 - Authored annotation guidelines and partnered with engineering and applied research to translate linguistic findings into ranking and recall improvements.

EDUCATION

- UC Davis

2022

Data Analytics / Data Science Bootcamp
- UC Santa Cruz

2021

B.A. Linguistics

PROJECTS

- Citi Bike Tableau Visualization — NYC trip data dashboard on Tableau Public. Full portfolio at [jonahrahn.github.io](#).

SKILLS

- BACKEND

API design · REST · database design · data pipelines · observability · auth
- LLM EVAL

Adversarial prompting · failure-mode taxonomy · red-teaming · rubric design
- LANGUAGES

Python (pandas, NumPy, Jupyter) · SQL · JavaScript
- DATA

BigQuery · Looker · Tableau · ETL · A/B testing · data warehousing
- ML / AI

PyTorch · TensorFlow · LangChain · prompt engineering · NLP
- LINGUISTICS

Syntax · computational linguistics · annotation taxonomy · multilingual evaluation
- INFRA

GCP · cloud infrastructure · MongoDB · PostgreSQL · Flask · React