

Backend engineer with hybrid expertise in LLM evaluation and computational linguistics. Currently building backend systems for a consumer health product at Oasis Health. Previously shipped adversarial testing, failure-mode taxonomies, and structured-evaluation pipelines at Persona Identities and Scale AI. I bring evaluation discipline to production code.

EXPERIENCE

Oasis Health · Backend Engineer

Jan 2026 – Present

New Haven, CT

- Shipped backend APIs powering Oasis's in-store product-scan flow, translating raw lab data into consumer-facing contaminant safety scores. *[Add: requests/day or P95 latency.]*
- Designed and owned the data-ingestion pipeline that backs Oasis's contaminant database, supporting **[N]** product categories and **[N]** new lab sources.
- Hardened the production stack — auth, rate limiting, observability — so the mobile clients can be trusted in front of real consumer purchase decisions.

Persona Identities · Senior Product Development Specialist

Oct 2025 – Nov 2025

- Delivered two end-to-end LLM evaluation projects for enterprise clients, including pipeline review, reporting, and post-mortem documentation.
- Implemented security fixes that mitigated incoming risk across **5+** high-priority enterprise clients, closing identity and reasoning vulnerabilities before deployment.
- Authored evaluation reports and failure-mode taxonomies that fed into the team's debugging and refinement work.

Scale AI · Operations Associate, Trust and Safety (Eval Engineer)

Apr 2024 – Jul 2025

- Built adversarial prompts and structured evaluations that surfaced harmful, biased, and unreliable model outputs for enterprise clients.
- Authored weekly reports analyzing drift, misalignment patterns, hallucinations, and behavioral risk using SQL and Python.
- Designed evaluation criteria, rubrics, and taxonomies that improved reviewer precision and reproducibility across high-volume projects.
- Partnered with ML teams to refine failure modes and update eval pipelines.

Scale AI (via Hire Art) · Generative AI Specialist → Queue Manager

Jan 2024 – Apr 2024

- Evaluated LLM outputs for safety, bias, hallucinations, and factual accuracy; ranked completions against gold standards.
- Diagnosed breakdowns in GenAI evaluations, escalated root causes, and coordinated fixes with engineering and QA.
- Quantified reviewer throughput, variance, and pipeline performance with Python and SQL.

Meta (via Magnit) · Data Labeling Analyst, Redteaming

Nov 2023 – Dec 2023

- Performed adversarial testing on production LLMs and delivered detailed failure-mode reports that fed into model robustness work.

- Boosted ML model accuracy by **30%** on Seeker Relevance through targeted linguistic annotation and preprocessing improvements.
- Authored annotation guidelines and reviewed labeler output that fed into the Seeker Relevance training set.
- Partnered with engineering and applied research to translate linguistic findings into ranking and recall improvements.

EDUCATION

UC Davis	2022
Data Analytics / Data Science Bootcamp	
UC Santa Cruz	2021
B.A. Linguistics	

PROJECTS

- **Citi Bike Tableau Visualization** — interactive NYC trip data dashboard on Tableau Public.
- *Several projects in progress; full list on the portfolio.*

INTERESTS

LLMs · language research · mountain biking · soccer · podcasts

SKILLS

BACKEND	API design · REST · database design · data pipelines · observability · auth
LLM EVAL	Adversarial prompting · failure-mode taxonomy · red-teaming · rubric design
LANGUAGES	Python (pandas, NumPy, Jupyter) · SQL · JavaScript
DATA	BigQuery · Looker · Tableau · ETL · A/B testing · data warehousing
ML / AI	PyTorch · TensorFlow · LangChain · prompt engineering · NLP
LINGUISTICS	Syntax · computational linguistics · annotation taxonomy · multilingual evaluation
INFRA	GCP · cloud infrastructure · MongoDB · PostgreSQL · Flask · React